

THE SYSTEM IS THE PRODUCT

AI AGENT OPERATING SYSTEM

**From prompt chains to controlled
autonomy.**

A field guide to choosing the smallest
architecture that can carry the work.

USER INTENT

objective · context · constraints



CONTROL PLANE

route · plan · approve · evaluate



EXECUTION PLANE

tools · memory · state · outputs

Autonomy is a system boundary.

The question is not how agentic the system feels. It is where control is designed to live.

MAKE THE BOUNDARY VISIBLE.

A system earns autonomy only when its work, control, and learning loops can be inspected by a human.

01	WORK	What must be done task · context · definition of done
02	CONTROL	Who decides what happens next route · plan · permissions · escalation
03	LEARNING	How the system gets better trace · eval · feedback · release gate

Choose architecture by problem shape, not by fashion.

Two questions predict most of the design: can the steps be known, and can the work split safely?



M
O
R
E
U
N
C
E
R
T
A
I
N

SINGLE AGENT

open-ended · tool use

ORCHESTRATED

specialists · synthesis

WORKFLOW

known steps · low variance

PARALLEL

independent branches

LOW VARIANCE

MORE PARALLEL →

HIGH



ANTI-PATTERN

**More
agents
does not
mean
more
system.**

Complexity is the cost.